

# Today

- Recap & questions from homework
  - Stan etc installs?
- Maybe an example to work on

# Reproducible workflows

- Reproducible analysis reports
  - not necessary but nice
  - supplementary material, data repositories
- Reports from .R files
- Any .R file can be knitted
- Can also include markup
  - e.g. 03\_9\_train\_ssq\_grid.R

# McElreath Ch8

- Learning goals:
- Understand and use **MCMC algorithms** to sample from the posterior distribution
- Recognize and fix bad sampling scenarios
- Use **HMC** implemented in Stan via R packages

# Main points McElreath Ch8

- **MCMC**: Monte Carlo Markov Chain
- Series of random numbers where each number **depends** on the previous one
- Sample less from low probability areas; more bang for your random buck
- Algorithms
  - Metropolis-Hastings
  - Gibbs sampling
  - **HMC**: Hamiltonian Monte Carlo
- We'll mostly use HMC

# MCMC algorithms

## Algorithm (general)

for many iterations

- propose new value for parameter

- calculate the probability of accepting the proposal:

$$P_{\text{accpt}} = \Pr(\text{proposal}) / \Pr(\text{current})$$

- accept proposal randomly with  $\text{Bern}(P_{\text{accept}})$

plot posterior distribution (histogram) of parameter values

where  $\Pr() = \text{prior} \times \text{likelihood}$

# Rosenbluth algorithm

## aka Metropolis-Hastings

### Algorithm (original)

for many iterations

- propose new value for parameter:

  - draw  $\text{Unif}(-\text{max\_d}, \text{max\_d})$

  - proposal = current parameter + draw

- calculate the probability of accepting the proposal:

  - $P_{\text{acpt}} = \text{Pr}(\text{proposal}) / \text{Pr}(\text{current})$

  - accept proposal randomly with  $\text{Bern}(P_{\text{accept}})$

plot posterior distribution (histogram) of parameter values

where  $\text{Pr}() = \text{prior} \times \text{likelihood}$

# MCMC algorithms

- Two other important algorithms
- Gibbs sampling
  - needs conjugate priors
  - prior such that the posterior is the same as the prior
  - e.g. norm prior x norm lik = norm posterior
  - beta prior x binom lik = beta posterior
- Hamiltonian Monte Carlo (HMC)

# MCMC algorithms

- Get an intuition for their behavior:
- <https://chi-feng.github.io/mcmc-demo/app.html#HamiltonianMC,standard>



# Bayesian tools

- Stan
  - Gelman group
  - Hamiltonian Monte Carlo
  - Betancourt (2017) A conceptual introduction to Hamiltonian Monte Carlo (<https://arxiv.org/abs/1701.02434A>)
  - open source
  - models with continuous parameters only
  - state of the art
- <http://mc-stan.org>

# Bayesian tools

- BUGS (Bayesian inference Using Gibbs Sampling) (and Metropolis-Hastings)
- <http://www.openbugs.info>
- Older, original standard tool for MCMC
- Exceedingly difficult to run on Mac
- Many newer tools are based on BUGS code style
- Lots of books and publications use BUGS
- Recommend: need to know about historically but don't use anymore

# Bayesian tools

- JAGS (Just Another Gibbs Sampler)
- <http://mcmc-jags.sourceforge.net/>
  - cross platform, open source
  - basically the same as BUGS
  - often faster
  - highly recommended for models that can't be fit in Stan (e.g. discrete parameters)
  - easy install
- Best to run from R
  - Install R2jags package (install from R)

# Bayesian tools

- Others:
- Nimble
  - somewhat common in ecology
- Julia: Turing package

# Main points McElreath Ch8

- Using **HMC** via **Stan** to fit models
- Now getting **posterior samples** from **HMC**
- Use **ulam** in rethinking to do HMC to follow examples
- Same syntax as sampost

# Main points McElreath Ch8

ulam or sampost

```
m1 <- ulam(  
  alist(  
    y ~ dnorm(mu, sigma),  
    mu <- a + b * x,  
    a ~ dnorm(0, 100),  
    b ~ dnorm(0, 10),  
    sigma ~ dcauchy(0, 2)  
  ),  
  data=d1)
```

# Main points McElreath Ch8

## ulam or sampost

```
m1 <- ulam(  
  alist(  
    y ~ dnorm(mu, sigma),  
    mu <- a + b * x,  
    a ~ dnorm(0, 100),  
    b ~ dnorm(0, 10),  
    sigma ~ dcauchy(0, 2)  
  ),  
  data=d1)
```

## rstanarm

```
m1 <- stan_glm(y ~ b, data=d1)
```

- **rstanarm** also uses Stan, HMC
  - Stan group's linear models package
  - syntax like lm
  - sensible default priors


# Main points McElreath Ch8

- Good choice of priors (**weakly informative**) can be helpful to tame model fit
  - e.g. **Half-Cauchy** instead of uniform
- MCMC **diagnostics** to judge convergence of fit
  - $\hat{r}$ ,  $n_{\text{eff}}$
  - plot chain traces ("time series")
- **Visualize** posteriors
  - histograms, pairs plot



# Chains



Random  
starts



3 chains

# Chains

Bad



Good





Fig 8.5

# Chains

Not converged

Converged

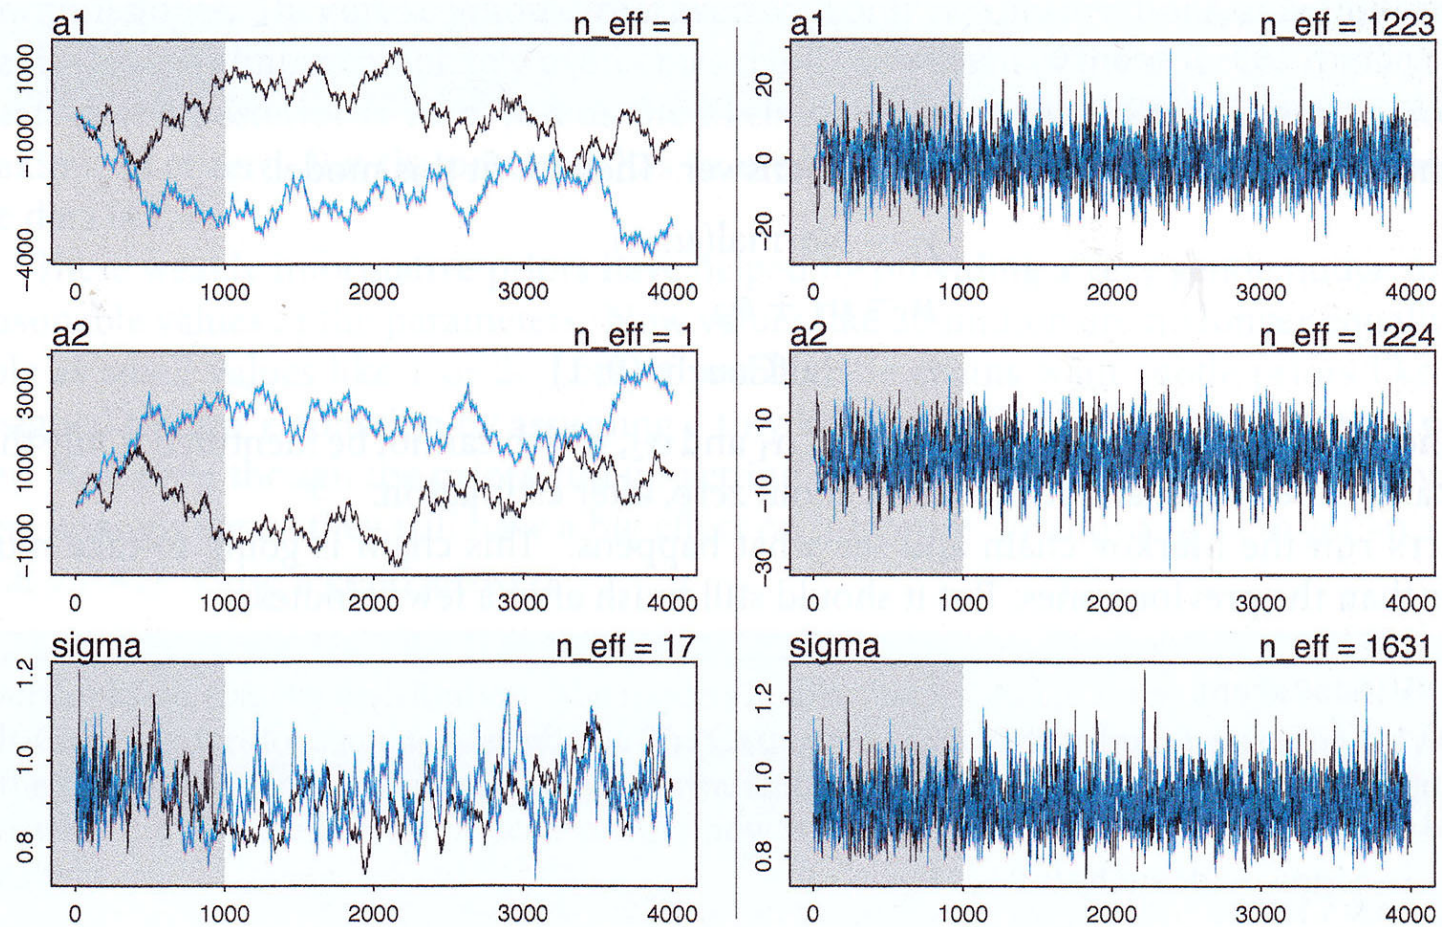


Fig 8.7

# How to fix

- Better starting values
- Weakly informative priors
- Uncorrelated parameters (e.g. standardized)
- Less common: adjust MCMC algorithm parameters